




list



Loupe: End-to-End Learning of Loop Unrolling Heuristics for Abstract Interpretation

Maykel Mattar^{1,2}, Michele Alberti¹, Valentin Perrelle¹, Salah Sadou²

1. *Université Paris-Saclay, CEA, List, Palaiseau, France*

2. *Université Bretagne Sud, IRISA, Vannes, France*

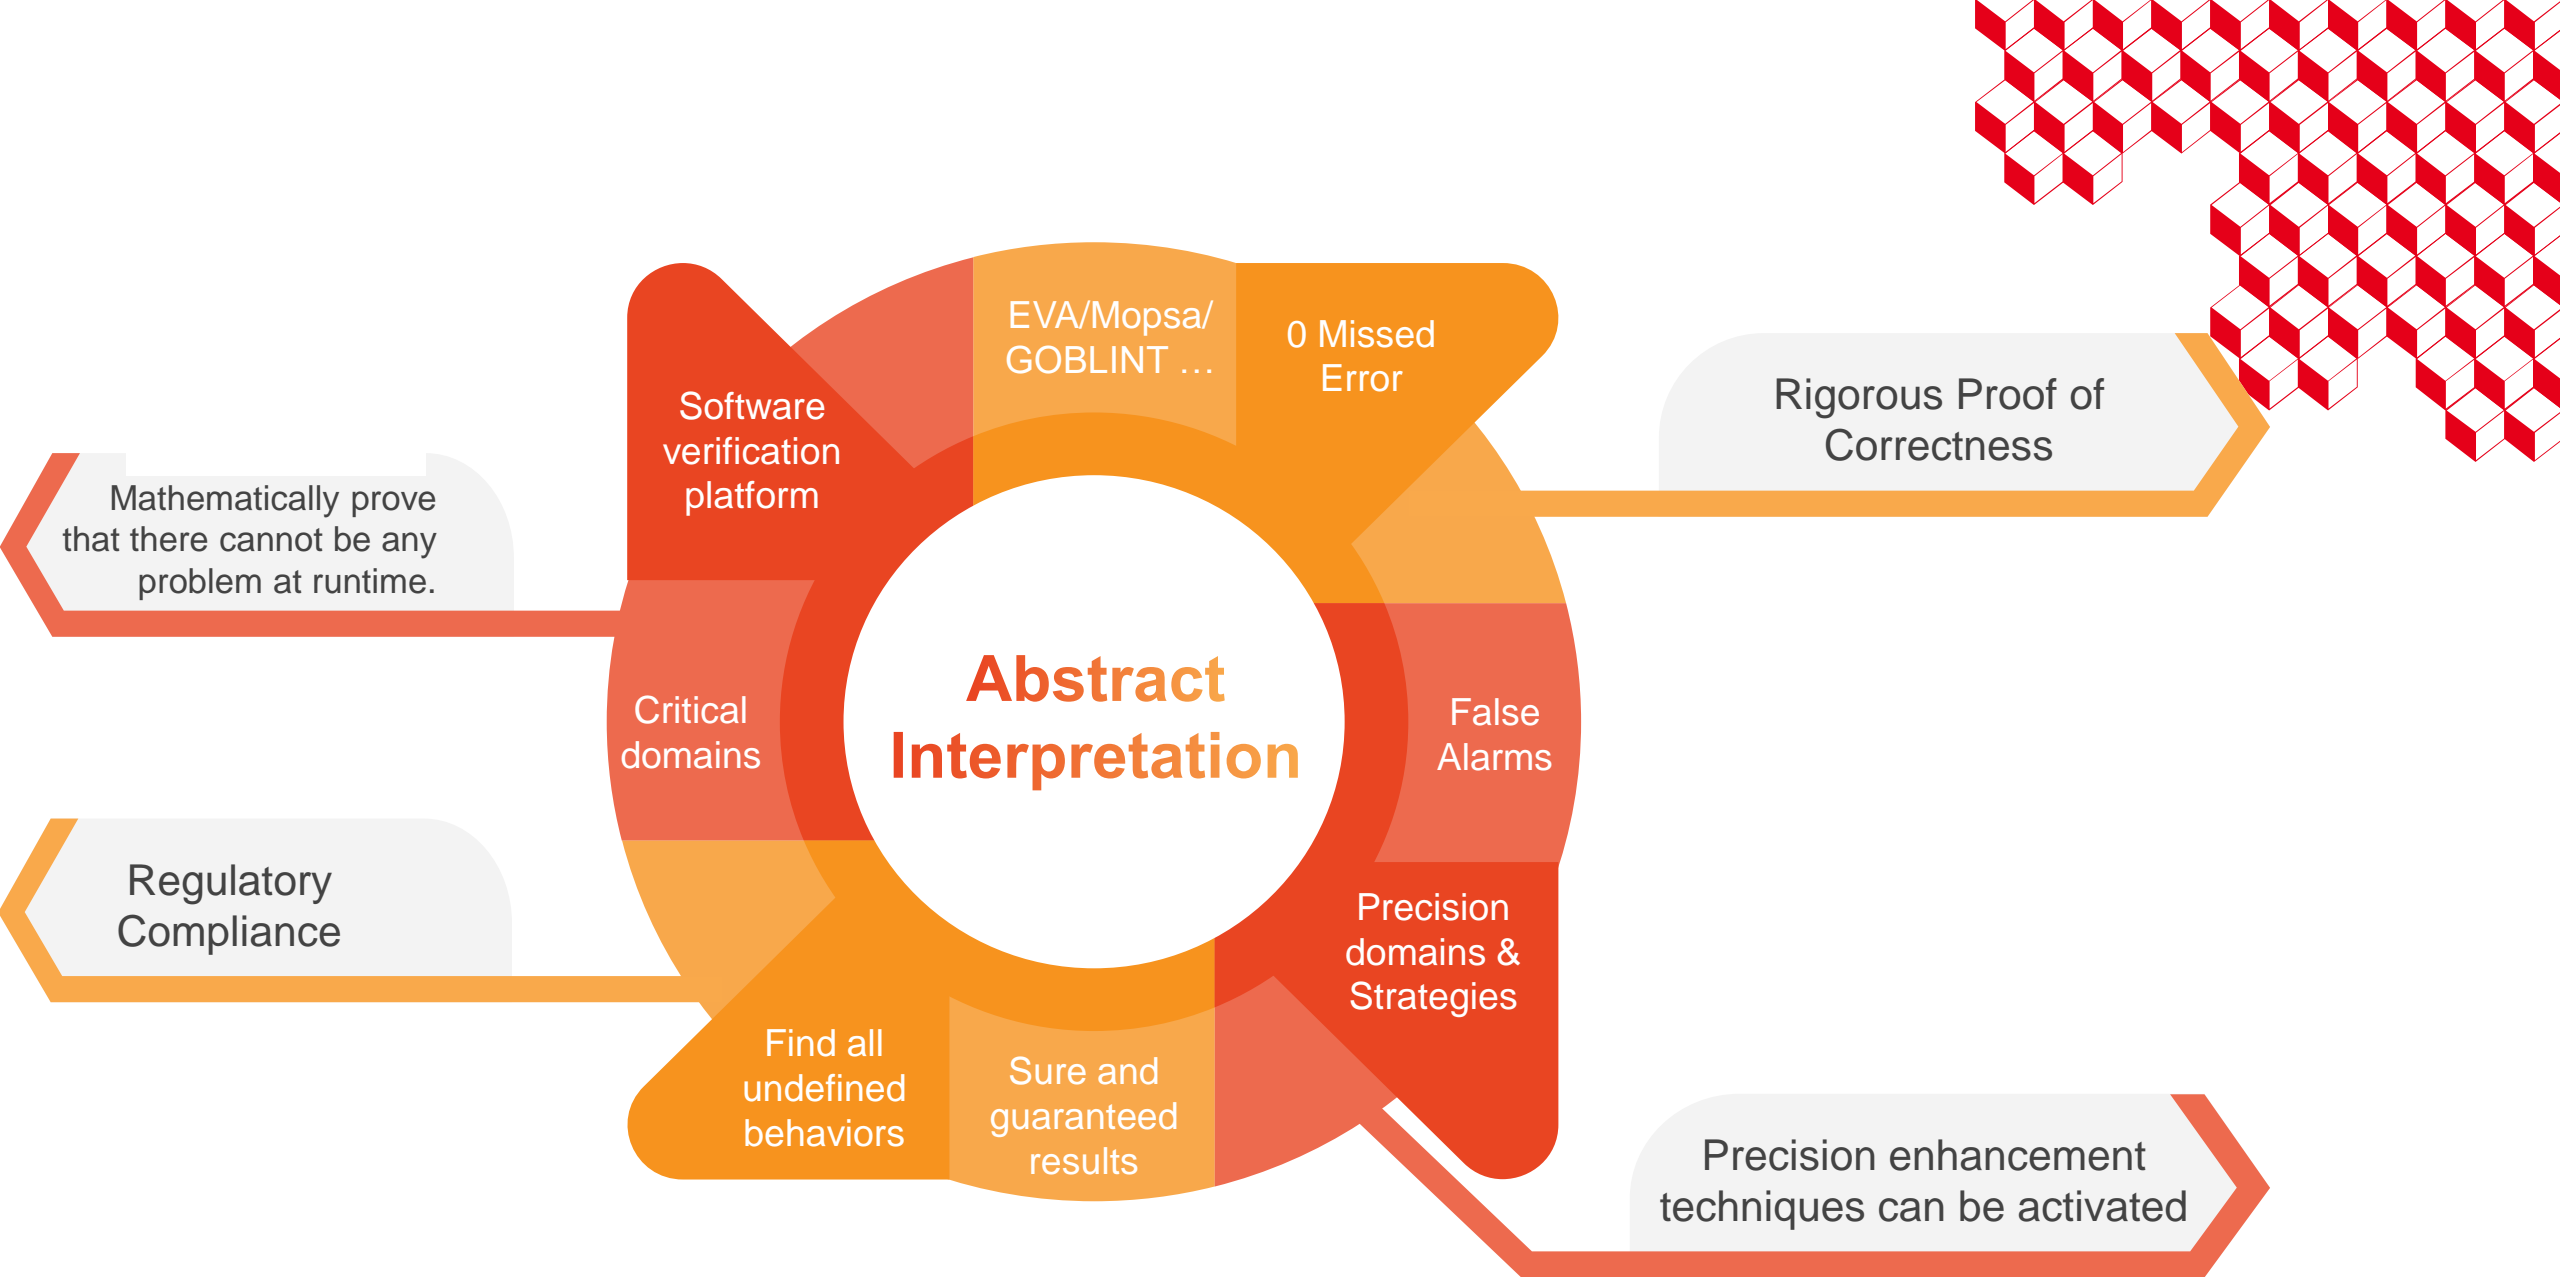


Agenda





1 ■ General context





**“ With great configurability
comes great complexity.**



Example

When Eva is **not** requested unroll the loop:



- **a** and **b** are in the range $[1, 2^{31} - 1]$
- Overflow alarm for the operation **a += b**.
- The interval abstraction of **a** and **b** fails to capture the relationship between these variables and **i**;
- Eva may not find a precise invariant before considering the entire positive range of **32-bit integers**

```
1.  #include <stdio.h>

2.  int fib(int n) {
3.      int a = 1, b = 1;
4.
5.      for (int i = 3; i <= n; i++) {
6.          int tmp = a;
7.          a += b;
8.          b = tmp;
9.      }
10.     return a;
11. }

12. void main() {
13.
14.     for (int i = 1, n; i <= 10; i++) {
15.         printf("Enter a number <= 30: ");
16.         scanf("%d", &n);
17.         if (n > 0 && n <= 30) {
18.             printf("fib(%d)=%d\n", n, fib(n));
19.             break;
20.         }
21.     }
22. }
```

Example

When Eva is requested to unroll the loop:



- The loop in line.4 is analyzed on iteration basis
- ***a and b*** range is precisely approximated.
- The alarms disappear


While loop unrolling can sometimes accelerate the analysis, it often increases computational cost significantly, particularly for nested loops.

```
1. #include <stdio.h>

2. int fib(int n) {
3.     int a = 1, b = 1;
4.     //@ loop unroll 100;
5.     for (int i = 3; i <= n; i++) {
6.         int tmp = a;
7.         a += b;
8.         b = tmp;
9.     }
10.    return a;
11. }

12. void main() {
13.     //@ loop unroll 0;
14.     for (int i = 1, n; i <= 10; i++) {
15.         printf("Enter a number <= 30: ");
16.         scanf("%d", &n);
17.         If (n > 0 && n <= 30) {
18.             printf("fib(%d)=%d\n", n, fib(n));
19.             break;
20.         }
21.     }
22. }
```

Parameterization Approaches



Approach Type	Representative Works / Tools	Key Idea / Technique	Limitations
Hard Coded Heuristics	<i>eva-auto-loop-unroll</i>	According to the user-provided factor, it attempts to unroll all loops that can be unrolled.	Factor-driven; over-unroll simple loops and overlook complex ones; Time consuming
Automated Tuning (Algorithmic Search)	PARF ¹ , TAILOR ²	Iterative execution with parameter optimization (local search, probabilistic refinement)	High computation cost; program-specific; no generalization
Machine Learning	Feature-based ^{3,4}	Learn heuristics from handcrafted program features	Requires feature engineering; not generalizable; expert bias
	Automatic Feature Generation – Sparrow ^{5,6}	Attempt to learn or generate features automatically	Limited to specific strategies; poor generality


1. Mansur et al., Automatically Tailoring Abstract Interpretation to Custom Usage Scenarios, CAV 2021.
2. Wang et al., Parf: Adaptive Parameter Refining for Abstract Interpretation, ASE 2024.
3. Oh et al., Learning a Strategy for Adapting a Program Analysis via Bayesian Optimisation, SIGPLAN Not., 2015.
4. Jeong et al., Data-Driven Context-Sensitivity for Points-to Analysis, Proc. ACM Program. Lang., OOPSLA 2017.

5. Jeon et al., Learning Graph-Based Heuristics for Pointer Analysis without Handcrafting Features, Proc. ACM Program. Lang., OOPSLA 2020.
6. Chae et al., Automatically Generating Features for Learning Program Analysis Heuristics for C-Like Languages, Proc. ACM Program. Lang., OOPSLA 2017.



3. Our Approach

Key Objectives



Use loop code;
End-to-End
learned heuristic



Auto-generate
labeled data

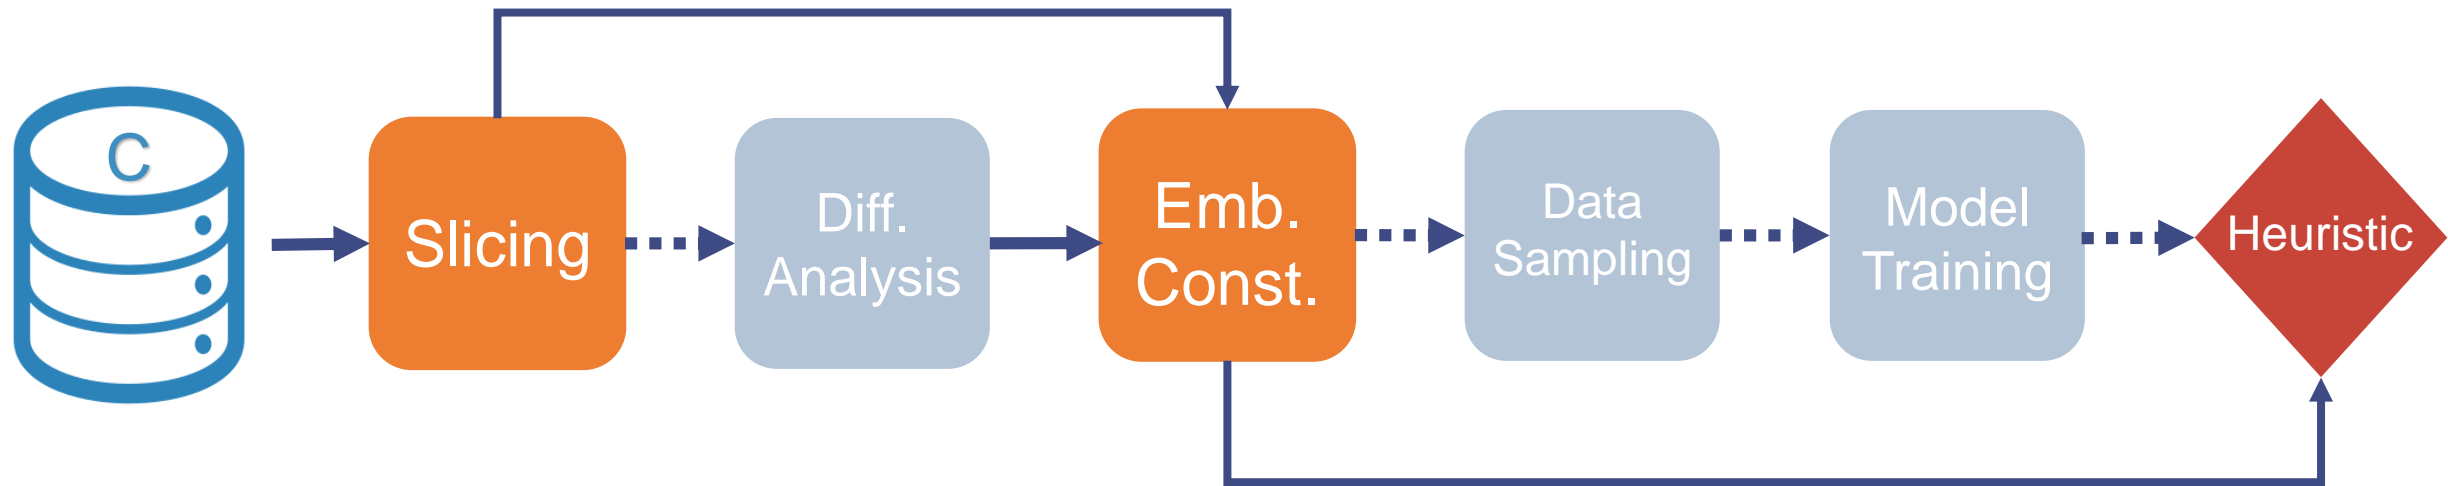


Select the
appropriate
learning methods



Keep the approach
modular and
adaptable

Loupe: End-to-End Learning of Loop Unrolling Heuristics for Abstract Interpretation



C Datasets

AnghaBench

TRAIN/EVAL/TEST


- Large-Scale Dataset >1M file
- Self-Contained Files
- Direct Processing
- Diverse Code Samples

Open Source Case Studies

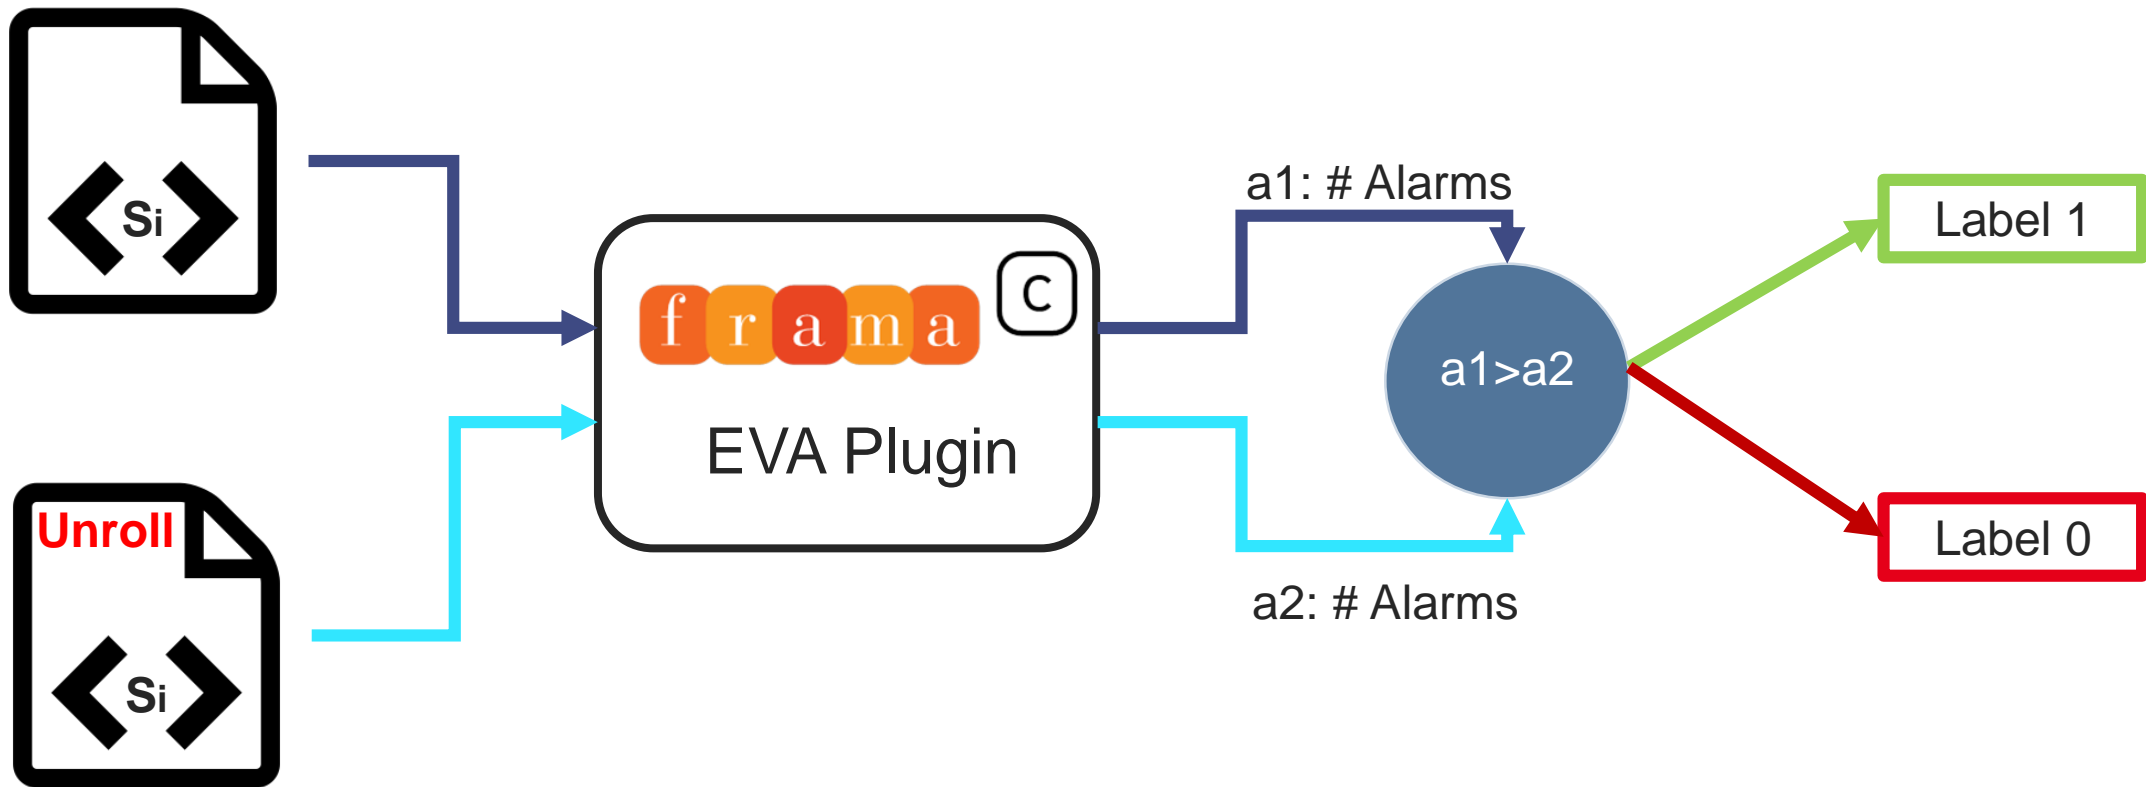
REAL-WORD AND PERFORMANCE EVALUATION

- Real-world C projects in their original form.
- Dependencies & complexity as found in actual software.
- Adapted for static analysis with Eva.
- Partially annotated.

Slicing






Differential Analysis

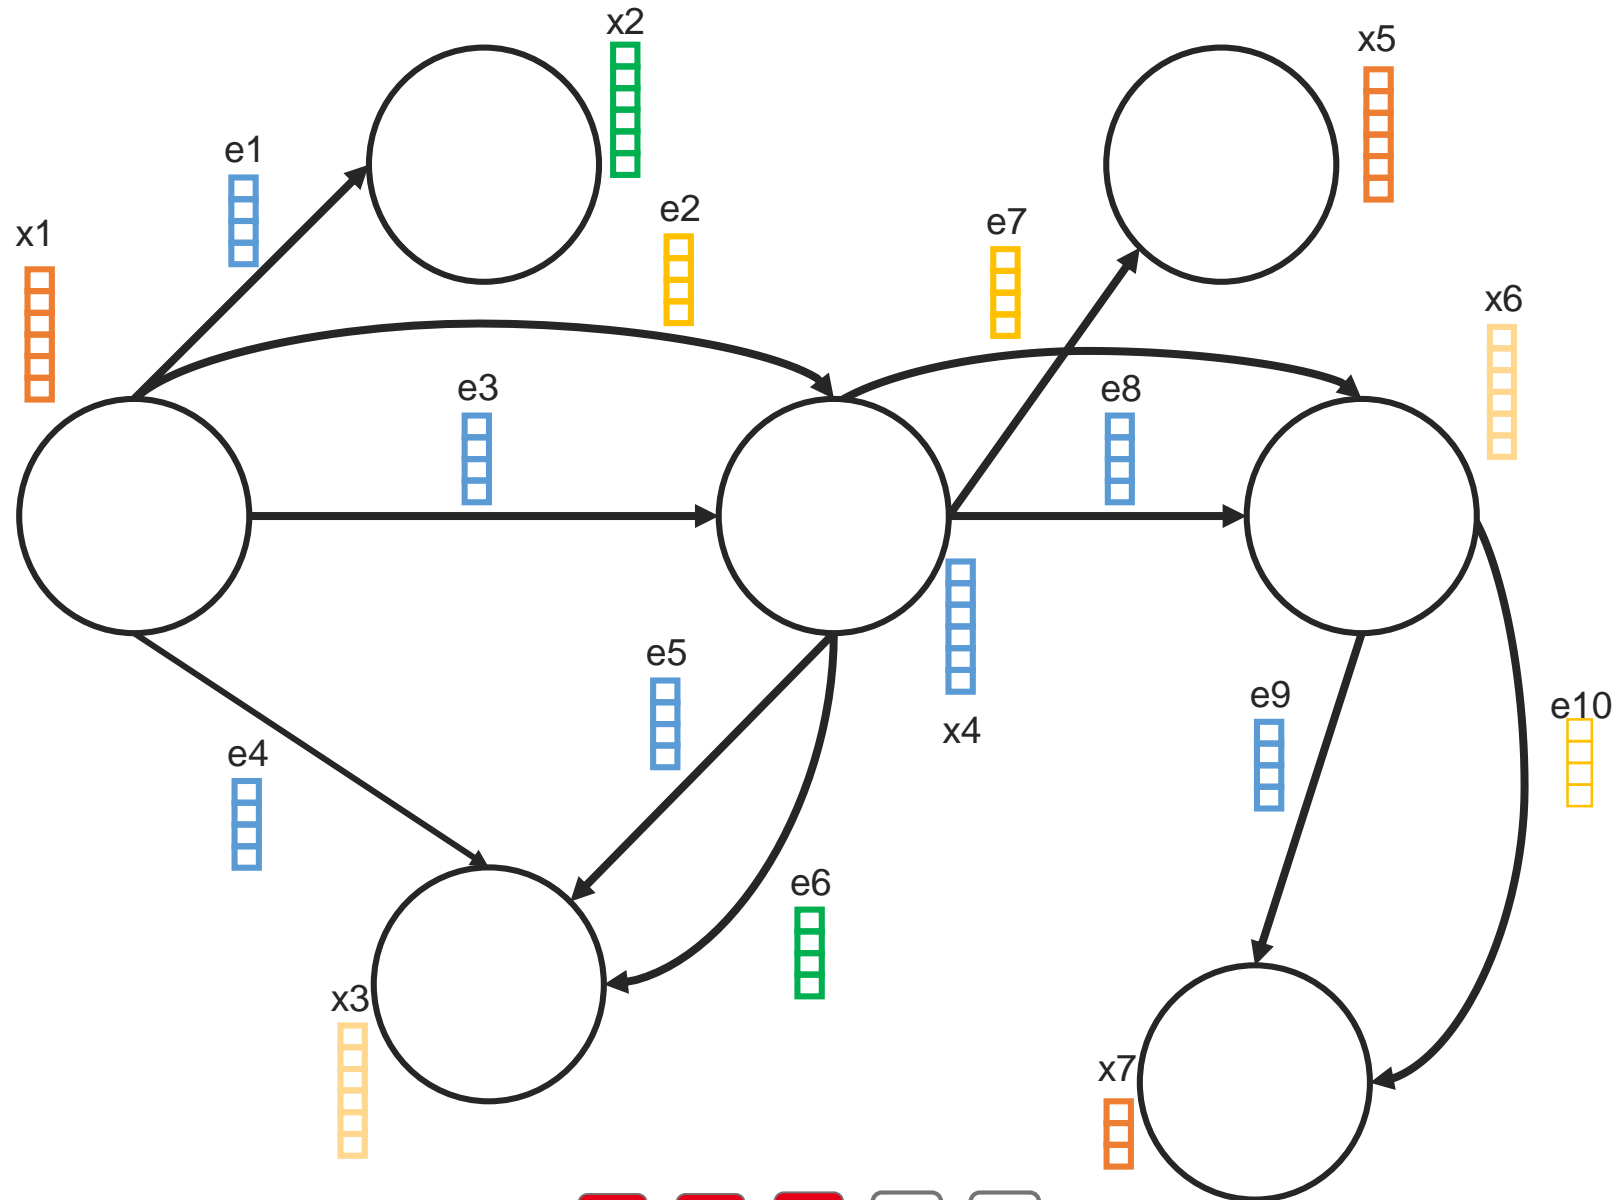


Embedding Construction

$$G_{\text{CPG}} = (V, E, \lambda, \mu)$$



$$G(s) = (X, \text{EdgeAttributes}, \text{EdgeIndex})$$



Data Sampling

Since loop unrolling is rarely required, the resulting dataset exhibits a significant imbalance, with ratio around 1:10 of positive to negative cases.

01

NAÏVE METHODS

- Undersampling
- Oversampling
- α -undersampling
- α -oversampling

02

DURING TRAINING

- Weight Balancing
- Focal Loss

03

GENERATION

- SMOTE
- SMOTE+TOMEK

Slicing


Diff.
Analysis

Emb.
Construct


Data
Sampling

Training


Model Training




Model Training




Model Training




Model Training



Model Training



Model Training





4. Results

Results



Family	Models Tested	Key Feature / Representation
Vector-based	XGBoost, SVM	Uses IR2Vec or handcrafted features (SPARROW) on linear code representations.
Code Language Models (CodeLLM)	GraphCodeBERT, CodeBERT	Transformer models pre-trained on code. Leverage pre-trained representations via fine-tuning.
Graph Neural Networks (GNN)	DGCNN, GAT, GIN, GINE	Programs represented CPG. Captures structural relationships in the code.



TABLE I: Performance comparison of loop unrolling heuristic learning on preprocessed and real-world imbalanced datasets

Model	ANGHABENCH (Test Split)				OSCS			
	P / R	F1	F2	B-ACC	P / R	F1	F2	B-ACC
Random	8.4 / 50.3	14.3	25.1	50.1	11.5 / 48.2	18.6	29.5	49.2
XGBOOST- IR2VEC	28.3 / 33.5	30.7	32.3	62.7	54.9 / 14.7	23.2	17.2	56.3
SVM- SPARROW	8.2 / 88.6	15.0	29.9	49.2	7.2 / 41.3	12.3	21.3	35.1
GRAPHCODEBERT	20.4 / 81.8	32.7	51.1	76.6	23.3 / 80.7	36.1	54.0	66.2
DGCNN	24.3 / 79.3	37.4	54.8	78.7	21.7 / 50.8	26.6	34.70	58.3
GAT	22.8 / 73.6	35.6	52.2	76.7	21.8 / 68.7	31.1	51.27	66.7
GIN	11.1 / 77.1	34.8	52.1	76.9	23.1 / 49.1	33.4	48.95	67.0
GINE	26.8 / 81.6	40.4	57.98	80.85	26.9 / 79.1	40.1	57.01	70.06

Performance Benchmark



	FRAMA-C (RQ2)							
Project	⊥		T		Alu-100		LOUPE (GINE)	
	#	Time	#	Time	#	Time	#	Time Alignment
ioccc								
mini-gmp								
genann								
papabench								
chrony								
kgflags								
libpng								
gnugo								
debie1								
basic-cwe-examples								
jsmn								
microstrain								
bench-moerman2018								
khash								
c-testsuite								
icpc								
solitaire								
line-following-robot								
safestringlib								
stmr								
powerwindow								
tweetnacl-usable								
qlz								
2048								
hiredis								
verisec								
c-utils								
Total**								

Performance Benchmark



	FRAMA-C (RQ2)								
Project	⊥		T		Alu-100		LOUPE (GINE)		
	#	Time	#	Time	#	Time	#	Time	Alignment
ioccc									
mini-gmp									
genann									
papabench									
chrony									
kgflags									
libspng									
gnugo									
debie1									
basic-cwe-examples									
jsmn									
microstrain									
bench-moerman2018									
khash									
c-testsuite									
icpc									
solitaire									
line-following-robot									
safestringlib									
stmr									
powerwindow									
tweetnacl-usable									
qlz									
2048									
hiredis									
verisec									
c-utils									
Total**									

⊥ (Analyzer with No Unrolling): The baseline configuration where no loops are unrolled.

T (Analyzer with All Unrolling): The configuration where all loops in the program are fully unrolled.

ALU-100: Eva built in heuristic *-eva-auto-loop-unroll* with 100 as a factor

Loupe (GINE): Our approach with the best performing model

Performance Benchmark



Project	FRAMA-C (RQ2)								
	⊥		T		Alu-100		LOUPE (GINE)		
	#	Time	#	Time	#	Time	#	Time	Alignment
ioccc									
mini-gmp									
genann									
papabench									
chrony									
kgflags									
libspng									
gnugo									
debie1									
basic-cwe-examples									
jsmn									
microstrain									
bench-moerman2018									
khash									
c-testsuite									
icpc									
solitaire									
line-following-robot									
safestringlib									
stmr									
powerwindow									
tweetnacl-usable									
qlz									
2048									
hiredis									
verisec									
c-utils									
Total**									

Alarms: The number of alarms reported by the analyzer (Lower is better/more precise).

Time (s): The total analysis time in seconds (Lower is better).

Alignment: Loupe (GINE) predictions vs. expert annotations. Higher ratio (closer to 1/1) is better.

- **TO:** Timeout of 2700s
- **Bold values:** Best performing result w.r.t. metric
- **Underlined values:** Better than ⊥, but worse than T
- *: Loupe (GINE) outperforms or matches Alu-100

Performance Benchmark




Project	FRAMA-C (RQ2)								
	⊥		T		Alu-100		LOUPE (GINE)		
	#	Time	#	Time	#	Time	#	Time	Alignment
ioccc	81	114.25	77	675.94	80	110.20	78*	127.56	1/1 27/38
mini-gmp	64	1.88	64	TO	64	2.02	64*	1.45	
genann	235	3.38	136	58.76	232	4.18	169*	4.80	
papabench	40	5.15	40	54.92	40	5.15	40*	3.68	
chrony		TO		TO		TO		TO	
kgflags	4	4.69	4	976.12	4	4.80	4*	102.05	4/4 5/6
libspng	237	11.17	233	44.59	237	11.58	237*	8.43	
gnugo	114	21.84	114	TO	114	TO	100*	27.22	
deb1el	31	39.54	9	1158.30	20	81.53	11*	110.77	
basic-cwe-examples	1	1.62	1	1.17	1	1.69	1*	0.70	
jsmn	68	20.83	2	12.54	68	21.14	2*	66.35	1/1 6/6
microstrain	1287	36.51	1287	TO	699	25.61	700	19.35	
bench-moerman2018	3	39.29	3	28.88	3	41.49	3*	34.74	
khash	1	0.13	0	0.05	1	0.12	1*	0.03	
c-testsuite	0	34.32	0	TO	0	TO	0*	23.49	
icpc	1	3.17	1	7.72	1	3.10	1*	2.30	2/2 5/5
solitaire	182	2.31	182	TO	173	4.17	173*	18.46	
line-following-robot	2	0.94	2	4.22	2	0.94	2*	0.64	
safestringlib	987	39.94	987	TO	845	69.19	453*	267.48	
stmr	67	5.21	67	602.19	67	5.38	67*	4.94	
powerwindow	1	18.18	1	67.33	1	18.17	1*	14.27	
tweetnacl-usable	99	3.26	3	26.32	4	25.54	56	4.96	
qlz	26	11.13	26	TO	26	39.02	26*	130.15	
2048	30	0.99	10	55.19	13	4.54	12*	1.18	
hiredis	248	110.69	210	244.97	224	101.75	241	92.39	
verisec									
c-utils									
Total**	3809	3230.488	3459	28319.22	2919	8681.39	2442*	3766.84	51/63 (81%)

- ✓ Matches approximately **67%** of the alarms reduced by T
- ✓ **7.5x** faster than T
- ✓ **5 TO** less than T
- ✓ Reduce around **35%** of alarms compared to ⊥
- ✓ Outperforms or matches Alu-100 **87%** of the time with almost half the time.
- ✓ Aligns with **51 out of 63** manually labeled unrolling decisions (81%)

Loupe (GINE) VS Alu



Let's examine how the learned heuristic Loupe (GINE) stacks up against the built-in heuristic (Alu) across varying unrolling factors.




- ✓ Increasing precision (low alarms) costs a massive, **non-linear increase** in time for Alu
- ✓ The learned model achieves **consistent efficiency** while maintaining high precision.
- ✓ Loupe (GINE) **selectively** unrolls required loops.
- ✓ Loupe (GINE) provides the **superior balance**, beating the built-in heuristic even at its extreme settings.

Consistency w.r.t Abstract Domains



Let's examine the consistency of our approach and how its alarm reduction trends perform across different abstract domain settings

- ✓ Alarm reduction trends remain **uniform** and predictable
- ✓ Loupe (GINE) **consistently outperforms** the built-in Alu heuristic in total alarm count on all configurations
- ✓ The absence of anomalies confirms the approach's **robustness and scalability** across varied configurations.



Transferability



MOPSA (RQ3)							
⊥		T		LOUPE (GINE)			
#	Time	#	Time	FRAMA-C #	Time	MOPSA #	Time
398	5.75	337	1877.12	<u>386</u>	338.83	<u>383</u>	489.98
15	1.60	15	1.82	15	7.90	15	1.67
5700	20.59	5700	58.65	5700	47.02	5700	43.74
3	0.04	3	0.14	3	0.03	3	0.04
42	1.18	40	2.22	40	1.00	40	1.00
64	0.36	61	1.12	64	0.25	64	0.24
392	21.32	213	18.04	<u>249</u>	16.21	<u>245</u>	15.98
19	0.34	19	TO	19	0.26	19	0.25
156	2.58	65	432.85	<u>117</u>	3.22	<u>112</u>	3.26
1	0.25	1	0.51	1	0.18	1	0.51
729	1.03	729	TO	692	194.98	692	567.24
55	0.42	55	46.74	55	0.32	55	0.35
11	0.28	11	0.62	11	0.26	11	0.21
142	3.68	68	9.16	<u>106</u>	3.78	68	8.76
1113	11.52	793	68.55	<u>974</u>	11.18	<u>853</u>	39.25
380	2.87	375	33.26	<u>377</u>	17.79	375	22.00
9220	73.88	8485	7950.86	<u>8809</u>	636.93	<u>8636</u>	1194.5

- ✓ Loupe (GINE) trained on Frama-C/EVA data **improved** performance when applied to MOPSA.
- ✓ Retraining Loupe (GINE) for MOPSA achieved **75%** of T's alarm reduction.
- ✓ The MOPSA-specific model was **6.5x** faster than the T strategy.



list



Merci

Maykel Mattar^{1,2}, Michele Alberti¹, Valentin Perrelle¹, Salah Sadou²

1. *Université Paris-Saclay, CEA, List, Palaiseau, France*

2. *Université Bretagne Sud, IRISA, Vannes, France*



maykel.mattar@cea.fr

maykel.mattar@univ-ubs.fr

